

Massively Parallel Analysis System for Seismology (MsPASS)

Prof. Gary L. Pavlis
Department of Earth and Atmospheric Sciences
Indiana University, Bloomington, IN

Overview

- Presentation on what MsPASS is and why you should use it
- Hands on session working through jupyter notebook



1. Massively Parallel

- Single CPU thing of the past
- Archaic software infrastructure

2. Analysis System for Seismology

- Research analysis - flexibility
- For Seismology research - tool for us not computer scientists

<https://www.mspass.org/>

Discussion

- MsPASS is a framework for research computing. What does the “research” keyword mean in this context?

Production versus Research IT systems

Production

- Solves Specific Problem
- Performance is critical
 - Time is money
 - Mission critical role with repeating tasks
- Operable with minimum skill set necessary for job
- Data model well known and fixed

Research

- Handle range of problems
- Performance is secondary
 - Feasible sufficient
 - Many one-up solutions
- Can assume users are specialists and life-long learners
- Data highly variable (one person's signal is another's noise)

Key Roles MsPASS can play in research today

- Manage large data sets
- Tools for low-level work to assemble event-based data
- Generic parallel tools for handling continuous data
- Complete set of tools for P-wave receiver function processing
- Tools for spectral estimation

Hands on exercise
focus

What is MsPASS good for

- Generic waveform processing toolkit
- Scalability
 - Prototype on a desktop
 - If needed scale up to HPC or cloud
 - Likely to become the tool of choice for interacting with Earthscope DMC when they get their cloud system working
- With new Earthscope system your MsPASS notebook will allow reproducing at least the low-level preprocessing to assemble your dataset

What is in the box?

- Exceptionally complete documentation (<https://mspass.org>) for open-source package
- Database to manage large data sets
- Parallel processing
 - HPC or cloud functionality
 - Clean error handling for long running jobs
- Flexible but simple to use IO abstractions (read-write almost any format and from URL)
- Python job control language
 - Reproducibility automatic if you submit jupyter notebook with your paper
- Algorithms - low level preprocessing (solved problem) focus
 - All obspy signal processing algorithms
 - Low-level windowing, bundling, and similar grungy stuff
 - Three-component primitives (rotation and transformation)
 - Trace editing
 - Header math
 - SNR calculations
 - All common P receiver function deconvolution+novel new method
- Docker container or anaconda package distribution

Sermonette/Sales Pitch

- This community needs to recognize what we have given you
- MsPASS makes archaic
 - SAC
 - AH
 - Most, if not all, existing generic waveform handling tools
- MsPASS supercedes ObsPy
 - If you use ObsPy you can use MsPASS with minimal effort
 - Drastically increases functionality over ObsPy

I know ObsPy well, why should I move to MsPASS?

- If you are only working with 10000 or less waveforms at a time it won't help
- For large data sets
 - Integrated database for data management essential
 - Clean error handling to run large jobs running for hours
 - Parallel processing functionality
 - Cross-platform compatibility - runs on desktop to highest end supercomputer or cloud system
- Likely to become the standard for low-level interaction with new Earthscope cloud system
- Reproducibility

Documentation, documentation, documentation:

- [User Manual](#)
- [Python API Reference](#)
- [C++ API](#)
- [Tutorials](#) - you all should have cloned this repository already

MsPASS is NOT

- Ideal solution for problems we viewed as solved:
 - Real time data handling
 - Seismic event catalog processing
 - Seismic reflection processing
 - Archival data management for Earthscope
- A complete solution for all of seismology
 - This is a framework, not a turnkey solution to all problems - may require some work to develop your custom “workflow”
 - More custom solutions require contributions from people like you

Hands on exercise part one: Data management

- Indexing miniseed files
- Managing source metadata
- Managing receiver (station) metadata
- Seismic data reader and writer abstraction

MsPASS Data Management Technology

- MongoDB is the integrated database of MsPASS
- MongoDB is a “document database”
 - Jargon name: a document == python dictionary
 - A “document” is the ultimate seismic header
 - You can save almost anything there
 - Completely flexible (by default) in what you save
- See the Data Management section of the MsPASS User Manual

To Jupyter Notebook

- Launch docker container as described [here](#)
- Probably will want to append “:v2.0.1” tag to run line

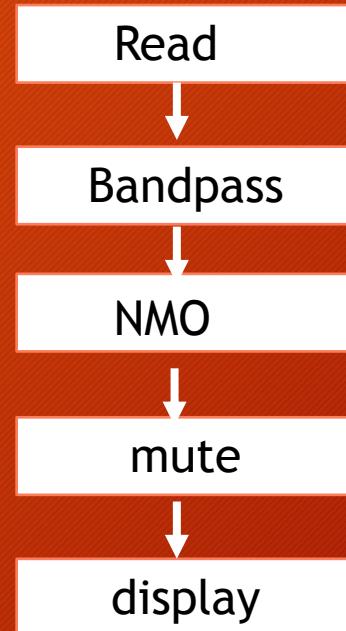
Part 2: Parallel schedulers

- Map-reduce model
 - Modern jargon term
 - Online sources obscure some simple concepts
 - Main paradigm for MsPASS parallel processing
- A modern paradigm likely to have a long lifetime
- I will introduce by analogy to unix pipeline concepts I assume all of you know

Seismic Reflection Workflow model

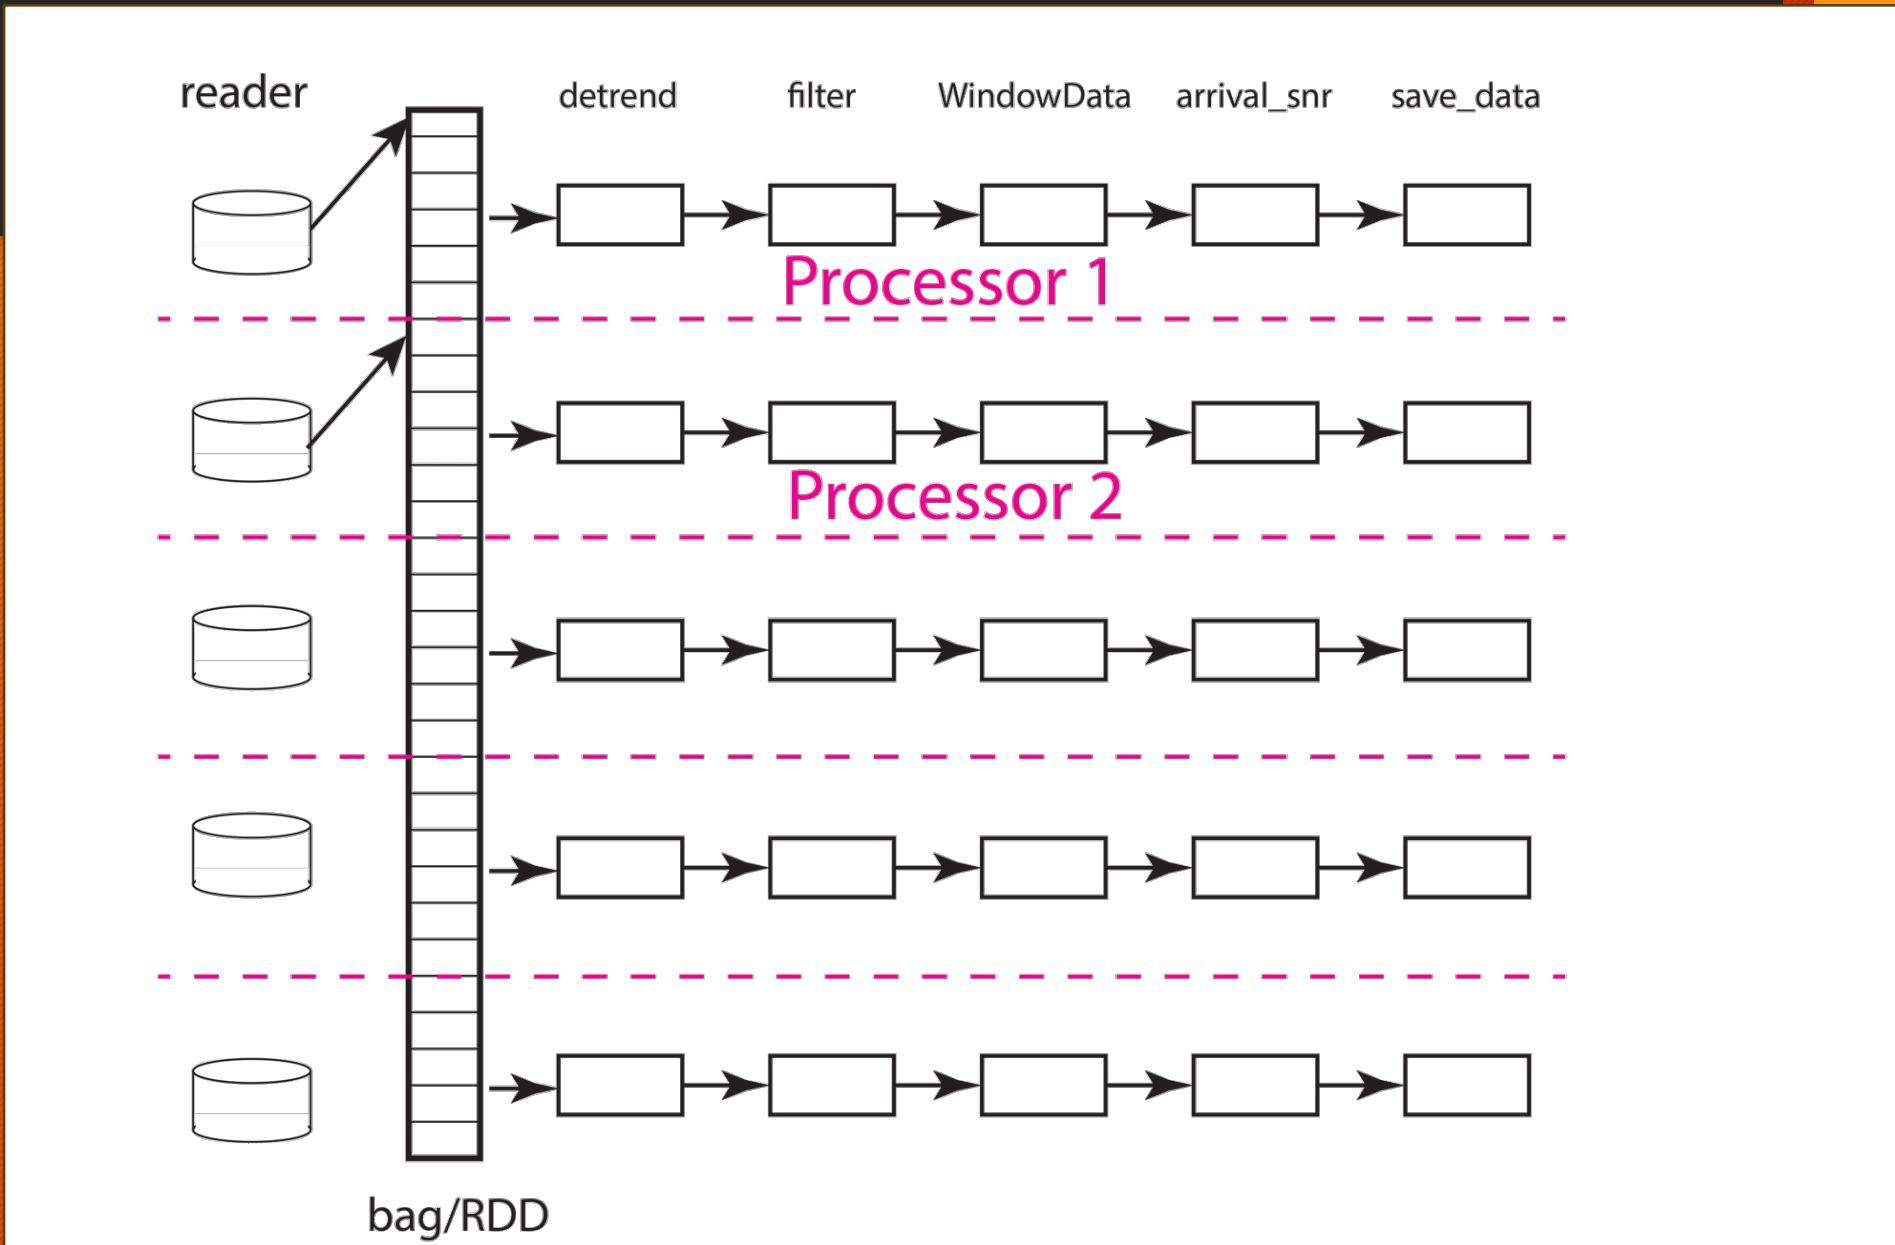
Example: seismic unix

- Illustrates traditional model
- Key point is data flow through processors
- Processors read input, modify it, and emit output

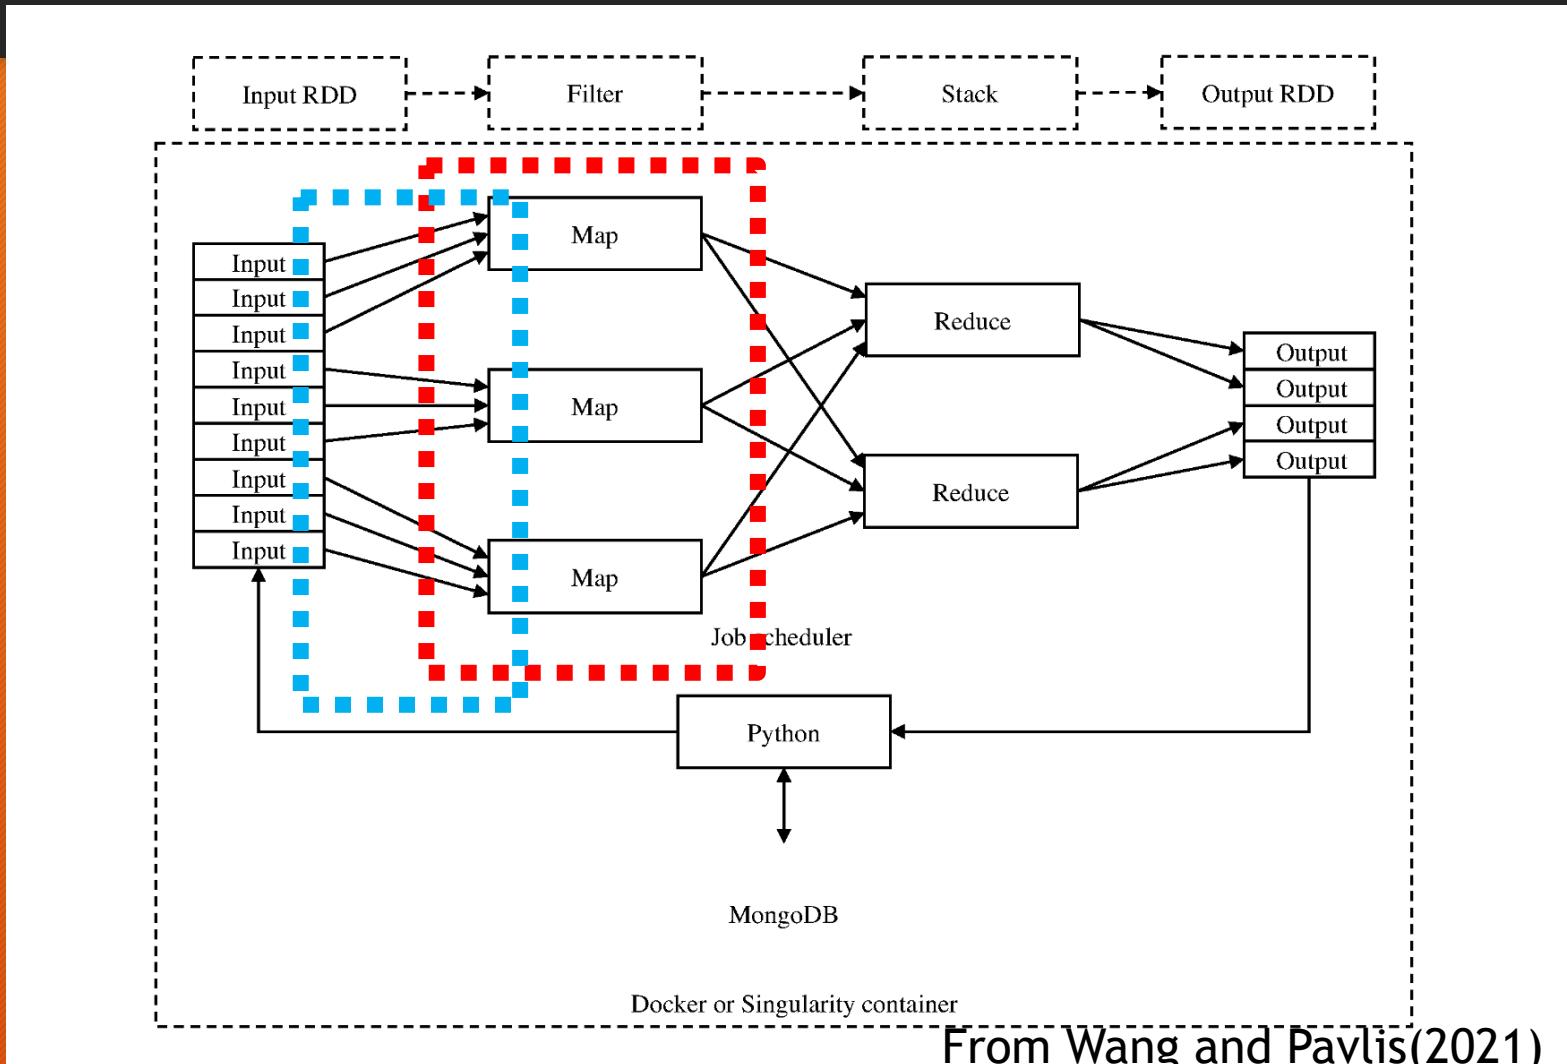


```
# real examples would add arguments for parameters  
subfilt < mydata | sunmo | sumute | suxwig
```

How to parallelize a similar job (2 processors)



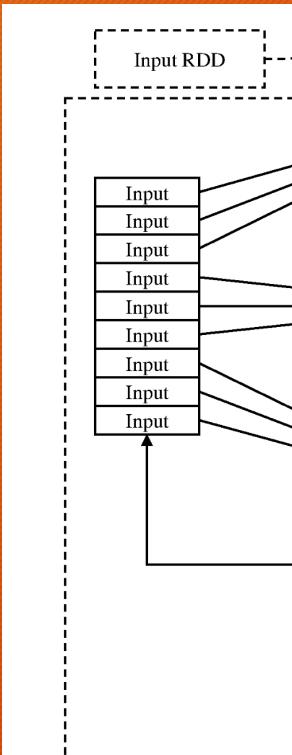
MsPASS Parallelization: map-reduce model



Key points

- **Map operator**
 - Behaves like a unix filter
 - Scheduler assigns each datum to processes
- **Reduce operator**
- **Data flow**
 - Parallel pipeline
 - Conceptually similar to unix shell | symbol
 - Scheduler moves data
 - Faster for threads than nodes

Spark RDD == Dask bag



- Documentation for both Spark and Dask obscure this topic
- RDD/bag concepts
 - Algorithmically identical to a large array of things (objects)
 - Workers can pull any component in equal time
 - The collection of things (data set) may not fit in memory
- Other parallel containers
 - Dataframe (table)
 - Array (matrix bigger than memory)

Comparison of serial and parallel workflows

Serial

```
cursor = db.wf_Seismogram.find({})
for doc in cursor:
    d = db.read_data(doc,collection="wf_Seismogram")
    d = signals.detrend(d,'demean')
    d = signals.filter(d,"bandbass",
        freqmin=0.01,freqmax=2.0)
    d = WindowData(d,200.0,500.0,t0shift=d.t0)
    db.save_data(d,collection="wf_Seismogram",
        data_tag="results")
```

Parallel (Dask)

```
cursor = db.wf_Seismogram.find({})
data = read_distributed_data(db, cursor)
data = data.map(signals.detrend,'demean')
data = data.map(signals.filter,"bandpass",
    freqmin=0.01, freqmax=2.0)
# windowing is relative to start time. 300 s window
# starting at d.t0+200
data = data.map(lambda d : WindowData(d, 200.0,
    500.0, t0shift=d.t0))
res = data.map(db.save_data,
    collection="wf_Seismogram",data_tag="results")
data_out = data.compute()
```

Key point: loop processing easily translated to series of map operators

Result acts like: detrend < datafile | filter > outfile